

# Deep Learning with Keras and Tensorflow (M3)

Created	@December 12, 2025 10:55 AM
Module Code	IBMM03: Deep Learning with Keras and Tensorflow

## Module 2: Advanced CNNs in Keras

### Advanced CNN Techniques Using Keras

#### What are CNNs?

Convolutional Neural Networks (CNNs) process visual data (Image).

What is an Image?

- A grid of numbers (Pixels)
- Each Pixel has 3 values: R, G, B
- Tensorflow represents this as a tensor:

Height × Width × Channels

e.g., 224 × 224 × 3

#### 👉 What are tensors? - Data Values and Shape

- mathematical object that generalizes scalars, vectors and matrices into a single concept.
- Tensors act as Multidimensional arrays / Data structure in ML and DL as Tensorflow stores data as tensors.
- Eg. Image: 3D Tensor (h, w, channels)
- A batch of Images: 4D Tensors
- A video: 5D Tensors

They use:

- **Convolution layers** → extract features
- **Pooling layers** → reduce size
- **Fully connected layers** → classification

A basic CNN stacks multiple Conv → Pool → Dense layers.

---

## Basic CNN

Typical structure:

1. Conv2D (32 filters, 3×3, ReLU)
2. MaxPooling (2×2)
3. Conv2D (64 filters)
4. MaxPooling
5. Flatten
6. Dense(128, ReLU)
7. Output Dense(10, Softmax)

Compiled with Adam + categorical crossentropy.

```
# Create a Sequential model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

Raw Image (jpg/png)

↓

Decode (convert to tensor)

↓

Resize to model's required size

↓

Normalize (0-1 or -1 to 1)

↓  
Apply augmentation (if training)  
↓  
Expand dims to batch shape  
↓  
Feed to CNN

---

## Advanced CNN Architectures

### VGG (Deep + Simple)

- Uses many stacked **3×3 convolution layers**
- Depth increases:  $64 \rightarrow 128 \rightarrow 256 \rightarrow 512$
- Followed by MaxPooling after each block
- Ends with fully connected layers

**Idea:** small filters + many layers → strong feature extraction.

```
model = Sequential([
    Conv2D(64, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(256, (3, 3), activation='relu'),
    Conv2D(256, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Dense(512, activation='relu'),
    Dense(512, activation='relu'),
    Dense(10, activation='softmax')
])
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
# Summary of the model
model.summary()
```

---

### ResNet (Residual Connections)

- Solves **vanishing gradient problem**
- Uses **skip connections**:
- Allow the network to learn identity mapping

$$\text{output} = F(x) + x$$

- Helps train very deep networks easily
- Residual block = Conv → Conv → Add shortcut → ReLU

**Key benefit:** stable training for deep models.

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization,
Activation, Add, Flatten, Dense

def residual_block(x, filters, kernel_size=3, stride=1):
    shortcut = x
    x = Conv2D(filters, kernel_size, strides=stride, padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Conv2D(filters, kernel_size, strides=1, padding='same')(x)
    x = BatchNormalization()(x)
    x = Add()([x, shortcut])
    x = Activation('relu')(x)
    return x

input = Input(shape=(64, 64, 3))
x = Conv2D(64, (7, 7), strides=2, padding='same')(input)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = residual_block(x, 64)
x = residual_block(x, 64)
x = Flatten()(x)
outputs = Dense(10, activation='softmax')(x)
```

## Data Augmentation Techniques

### What is Data Augmentation?

Data augmentation creates modified versions of training images (rotated, flipped, shifted, etc.) to make the model more robust and reduce overfitting.

### Why is it important?

- Increases data variety
- Prevents overfitting
- Improves generalization on unseen data

## Basic Augmentation Techniques

Using **ImageDataGenerator**, you can apply:

- **Rotation**
- **Width/height shifts**
- **Zoom**
- **Shear**
- **Horizontal flip**
- **Rescale**

These create new image variations during training.

```
from tensorflow.keras.preprocessing.image
import ImageDataGenerator

# Create an instance of ImageDataGenerator with augmentation options
datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Load a sample image and reshape it
from tensorflow.keras.preprocessing
import image
img = image.load_img('sample.jpg')
x = image.img_to_array(img)
x = x.reshape((1,) + x.shape)
```

## Advanced Augmentation

### Feature-wise normalization

- Dataset mean = 0
- Dataset std = 1

### Sample-wise normalization

- Each image mean = 0
- Each image std = 1

```
# Create an instance of ImageDataGenerator with advanced options
datagen = ImageDataGenerator(
    featurewise_center=True,
    featurewise_std_normalization=True,
    samplewise_center=True,
    samplewise_std_normalization=True
)
# Compute the mean and standard deviation on a dataset of images
datagen.fit(training_images)
# Generate batches of normalized images
i = 0
for batch in datagen.flow(training_images, batch_size=32):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break
plt.show()
```

## Custom Augmentation

You can write your own preprocessing function (e.g., adding random noise) and pass it to:

```
ImageDataGenerator(preprocessing_function=your_function)
```

```
import numpy as np
def add_random_noise(image):
    noise = np.random.normal(0, 0.1, image.shape)
    return image + noise
# Create an instance of ImageDataGenerator with custom augmentation
datagen = ImageDataGenerator(preprocessing_function=add_random_noise)
# Generate batches of augmented images
i = 0
for batch in datagen.flow(training_images, batch_size=32):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break
plt.show()
```

## Why do CNNs overfit?

Because they memorize patterns from LIMITED data.

So we trick the model by making **fake variations** of the same image.

TensorFlow can rotate, zoom, flip, crop, etc.

### Intuition:

Model sees:

- The same cat rotated
- The same cat zoomed
- The same cat with brightness changed

It learns:

👉 “*This is still a cat. I shouldn’t overfit to exact pixels.*”

Mathematically, all these transformations are just **function operations on pixel coordinates**, but you don’t need formula-level details right now.

## Transfer Learning in Keras

### What is Transfer Learning?

Using a model **pre-trained on a large dataset** (e.g., ImageNet) and reusing its learned features for a **new but related task**.

---

## Why use Transfer Learning?

- **Reduced training time** – model already knows useful features.
  - **Better performance** – pre-trained on huge datasets.
  - **Works well with small datasets** – especially useful in medical imaging, NLP, etc.
  - **Less computation needed.**
- 

## How it works

1. Load a **pre-trained model** (e.g., VGG16).
  2. Remove its top layers (`include_top=False`).
  3. Freeze its convolutional layers (`layer.trainable = False`).
  4. Add new custom layers for your task (Flatten + Dense layers).
  5. Compile and train.
- 

## Code Breakdown

### 1. Load Pre-trained Model

```
base_model = VGG16(include_top=False, input_shape=(224,224,3))
```

### 2. Freeze layers

```
for layer in base_model.layers:  
    layer.trainable = False
```

### 3. Add new classifier

```
model = Sequential()  
model.add(base_model)  
model.add(Flatten())  
model.add(Dense(256, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

### 4. Compile

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

### 5. Data preprocessing

```
train_datagen = ImageDataGenerator(rescale=1./255)  
train_generator = train_datagen.flow_from_directory(  
    directory,  
    target_size=(224,224),  
    batch_size=32,  
    class_mode='binary'  
)
```

### 6. Train

```
model.fit(train_generator, epochs=10)
```

## Fine-tuning

Unfreeze top layers to adapt the model more deeply:

```
for layer in base_model.layers[-4:]:
    layer.trainable = True
```

Then recompile + train again.

## Using Pre-Trained Models

### What is a Pretrained Model?

A **pretrained model** is a neural network that has already been trained on a large dataset (like **ImageNet**, with 14M images / 1000 classes).

Examples:

- **VGG16**
- **ResNet**
- **MobileNet**
- **EfficientNet**

These models have already learned:

- Edges
- Textures
- Shapes
- Object parts

...so you don't need to train them from scratch.

---

## Two Ways to Use Pretrained Models

### 1. Use as a Fixed Feature Extractor

- You **freeze** all layers → no retraining.

- Pass new images through the network.
- The model outputs **feature maps** (high-level patterns).
- You use these features for:
  - Clustering
  - Visualization
  - Dimensionality reduction
  - Feeding into ML models (SVM, RF, etc.)

## Benefits

- ✓ No training needed
- ✓ Very fast
- ✓ Works well with small datasets
- ✓ Uses very powerful learned features

## Limitations

- ✗ Cannot adapt to your dataset
- ✗ Performance limited if your data is very different from ImageNet

```

import os
import shutil
from PIL import Image
import numpy as np

# Define the base directory for sample data
base_dir = 'sample_data'
class1_dir = os.path.join(base_dir, 'class1')
class2_dir = os.path.join(base_dir, 'class2')

# Create directories for two classes
os.makedirs(class1_dir, exist_ok=True)
os.makedirs(class2_dir, exist_ok=True)

# Function to generate and save random images
def generate_random_images(save_dir, num_images):
    for i in range(num_images):
        # Generate a random RGB image of size 224x224

```

### How to use pretrained models as feature extractors in Keras

```

# Generate a random RGB image of size 224x224
img = Image.fromarray(np.uint8(np.random.rand(224, 224, 3) * 255))
# Save the image to the specified directory
img.save(os.path.join(save_dir, f'image_{i}.jpg'))

# Number of images to generate for each class
num_images_per_class = 100 # You can increase this to have more training data

# Generate random images for class 1 and class 2
generate_random_images(class1_dir, num_images_per_class)
generate_random_images(class2_dir, num_images_per_class)

print(f'Sample data generated at {base_dir} with {num_images_per_class} images per class.')

```

### Example: Pretrained model for extracting features

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam

# Load the VGG16 model pre-trained on ImageNet
base_model = VGG16(weights='imagenet'
, include_top=False, input_shape=(224, 224, 3))

# Freeze all layers initially
for layer in base_model.layers:
    layer.trainable = False

# Create a new model and add the base model and new layers
model = Sequential([
    base_model,
    Flatten(),
```

### Example: Pretrained model for extracting features

```
# Create a new model and add the base model and new layers
model = Sequential([
    base_model,
    Flatten(),
    Dense(256, activation='relu'),
    Dense(1, activation='sigmoid')

# Change to the number of classes you have
])

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001),
loss='binary_crossentropy', metrics=['accuracy'])
```

### Example: Pretrained model for extracting features

```
# Load and preprocess the dataset
train_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
    '/content/sample_data',
    target_size=(224, 224),
    batch_size=32,
    class_mode='binary'
)

# Train the model with frozen layers
model.fit(train_generator, epochs=10)

# Gradually unfreeze layers and fine-tune
for layer in base_model.layers[-4:]: # Unfreeze the last 4 layers
    layer.trainable = True

# Compile the model again with a lower learning rate for fine-tuning
model.compile(optimizer=Adam(learning_rate=0.0001),
loss='binary_crossentropy', metrics=['accuracy'])

# Fine-tune the model
model.fit(train_generator, epochs=10)
```

## 2. Fine-Tuning (Transfer Learning)

- You **unfreeze some top layers** of the pretrained model.
- Add new layers for your task.
- Train everything together (usually with low learning rate).

## Why Fine-Tune?

Because the pretrained model may not perfectly match your new dataset.

For example:

- ImageNet → dogs, cars, buildings
- Your dataset → retinal fundus images

So fine-tuning helps the model adjust.

## Benefits

- ✓ Much better accuracy
- ✓ Model adapts to your domain
- ✓ Still requires *far less* data than training from scratch

## Keras Workflow (Conceptual Steps)

### Step 1: Load pretrained model

```
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(2  
24, 224, 3))
```

### Step 2: Freeze its layers

```
base_model.trainable = False
```

### Step 3: Add your own classifier

Flatten → Dense(256, relu) → Dense(1, sigmoid)

### Step 4: Train your model (feature extraction mode)

This only trains the new layers you added.

### Step 5: (Optional) Fine-tune

Unfreeze top layers from VGG16 and train again with smaller LR.

## TensorFlow for Image Processing

Image processing means **transforming or analyzing images** to extract useful information.

Used in:

- Medical imaging
  - Autonomous vehicles
  - Facial recognition
  - Computer vision tasks
- 

## Why Use TensorFlow for Image Processing?

- **Easy to use** — high-level APIs simplify complex operations
  - **Pretrained models available** (ResNet, Inception, etc.)
  - **Scalable** from mobile devices to cloud clusters
  - **Strong community & documentation**
- 

## Basic Image Processing in TensorFlow

Common steps:

1. **Load image**
2. **Resize** (e.g., 224×224)
3. **Convert to NumPy or tensor**
4. **Add batch dimension** → required for model predictions

```
img = tf.expand_dims(img, axis=0)
```

```

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator,
img_to_array, load_img

# Load and preprocess the image
img = load_img('/content/path_to_image.jpg', target_size=(224, 224))
img_array = img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Add batch dimension

# Display the original image
import matplotlib.pyplot as plt
plt.imshow(img)
plt.show()

```

## Data Augmentation with TensorFlow

Used to increase dataset diversity → reduces overfitting.

Typical augmentations:

- Rotation
- Shifts (width/height)
- Shear
- Zoom
- Flip

TensorFlow can generate **batches of augmented images automatically**.

## Transpose Convolution

### What is Transpose Convolution?

Transpose convolution (also called **deconvolution**) is a technique used to **increase the spatial resolution** of feature maps.

It performs the **inverse of standard convolution** → used for **up-sampling**.

### Why Standard Convolution Is Not Enough?

- Normal convolution **reduces** spatial size.

- Tasks like **image generation**, **super-resolution**, and **semantic segmentation** need **larger** output images.
  - Transpose convolution expands the image.
- 

## How Transpose Convolution Works?

- Inserts **zeros between pixels** of the input feature map.
  - Applies a convolution filter over this expanded grid.
  - Produces a **larger** (up-sampled) feature map.
- 

## Applications

- **GANs**: generate images from a latent vector
  - **Super-resolution**: enhance image resolution
  - **Semantic segmentation**: create pixel-wise output maps
- 

## Keras Implementation

- `Conv2DTranspose(filters=32, kernel_size=3, strides=2, activation='relu')`  
→ upsamples by factor of 2
- Output layer often uses:
  - `sigmoid` (binary masks)
  - `softmax` (multi-class segmentation)
- Compile with:
  - Optimizer: `Adam`
  - Loss: e.g., `mse`, or segmentation-specific loss

```

import os
import logging
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2DTranspose

# Set environment variables to suppress TensorFlow warnings
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3' # Ignore INFO, WARNING, and ERROR
messages
os.environ['TF_ENABLE_ONEDNN_OPTS'] = '0' # Turn off oneDNN custom
operations

# Use logging to suppress TensorFlow warnings
logging.getLogger('tensorflow').setLevel(logging.ERROR)

# Define the input layer
input_layer = Input(shape=(28, 28, 1))

# Add a transpose convolution layer
transpose_conv_layer = Conv2DTranspose(filters=32, kernel_size=(3, 3),
strides=(2, 2), padding='same', activation='relu')(input_layer)

# Define the output layer
output_layer = Conv2DTranspose(filters=1, kernel_size=(3, 3),
activation='sigmoid', padding='same')(transpose_conv_layer)

# Create the model
model = Model(inputs=input_layer, outputs=output_layer)

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error',
metrics=['accuracy'])

# Display the model summary
model.summary()

```

## Issue: Checkerboard Artifacts

Caused by uneven kernel overlap during transpose convolution.

### How to Fix

Use:

- **UpSampling2D** → simple nearest-neighbor/bilinear upsampling
- How to fill pixel values in new positions:

#### a) Nearest Neighbor (most intuitive)

"Pick the closest pixel."

Works like zooming into a Minecraft block → blocky but fast.

Math (simple):

`new_pixel = old_pixel at nearest coordinate`

#### b) Bilinear (smoother)

"Blend the 4 nearest pixels."

Imagine mixing colors from nearby areas.

Math (concept version):

TensorFlow takes:

- top-left pixel
- top-right pixel
- bottom-left pixel
- bottom-right pixel

Then blends them depending on distance

- Followed by **Conv2D** → smooths & refines output

```
# Define a model with up-sampling followed by convolution to avoid
# checkerboard artifacts
x = UpSampling2D(size=(2, 2))(input_layer)
output_layer = Conv2D(filters=64, kernel_size=(3, 3), padding='same')(x)
```

This produces cleaner, artifact-free images.

Glossary :

Term	Definition
<b>Activation function</b>	A mathematical function used in neural networks to determine the output of a neuron.
<b>Adam optimizer</b>	An optimization algorithm that can be used instead of the classical stochastic gradient descent procedure to update network weights iteratively based on training data.
<b>Augmentation</b>	A process of increasing the diversity of training data by applying various transformations like rotation, scaling, and so on.
<b>Binary cross-entropy</b>	A loss function used for binary classification tasks, measuring the performance of a classification model whose output is a probability value between 0 and 1.
<b>Convolution</b>	A mathematical operation used in deep learning, especially in convolutional neural networks (CNNs), for filtering data.

Term	Definition
<b>Custom augmentation function</b>	A user-defined function that applies specific transformations to images during data augmentation, providing full control over the augmentation process.
<b>Data augmentation</b>	Techniques used to increase the diversity of training data by applying random transformations such as rotation, translation, flipping, scaling, and adding noise.
<b>Deconvolution</b>	Also known as transpose convolution, this is a technique used to up-sample an image, often used in generative models.
<b>Dense layer</b>	A fully connected neural network layer, where each input node is connected to each output node, commonly used in the final stages of a network.
<b>Feature map</b>	A set of features generated by applying a convolution operation to an image or data input.
<b>Feature-wise normalization</b>	A technique to set the mean of the data set to 0 and normalize it to have a standard deviation of 1.
<b>Fine-tuning</b>	The process of unfreezing some of the top layers of a pre-trained model base and jointly training both the newly added layers and the base layers for a specific task.
<b>Flatten layer</b>	A layer that converts the output of a convolutional layer to a 1D array, allowing it to be passed to a fully connected layer.
<b>Generative adversarial networks (GANs)</b>	A class of machine learning frameworks where two neural networks compete with each other to create realistic data samples.
<b>Height shift range</b>	A data augmentation parameter that randomly shifts an image vertically, altering its position to improve model robustness to vertical translations.
<b>TensorFlow Hub</b>	A repository of reusable machine learning modules, which can be easily integrated into TensorFlow applications to accelerate development.
<b>TensorFlow.js</b>	A library for training and deploying machine learning models in JavaScript environments, such as web browsers and Node.js.

Term	Definition
<b>Horizontal flip</b>	A data augmentation technique where the image is flipped horizontally, creating a mirror image to increase data diversity.
<b>ImageDataGenerator</b>	A Keras class used for generating batches of tensor image data with real-time data augmentation.
<b>ImageNet</b>	A large visual database designed for use in visual object recognition software research, often used as a data set for pre-training convolutional neural networks.
<b>Image processing</b>	The manipulation of an image to improve its quality or extract information from it.
<b>Kernel</b>	A small matrix used in convolution operations to detect features such as edges in images.
<b>Latent vector</b>	A vector representing compressed data in a lower-dimensional space, often used in generative models.
<b>Pre-trained model</b>	A model previously trained on a large data set, which can be used as a starting point for training on a new, related task.
<b>Random noise</b>	A type of custom augmentation that adds random noise to images, simulating different lighting conditions and sensor noise to make models more robust.
<b>Rotation range</b>	A data augmentation parameter that randomly rotates an image within a specified range of degrees, enhancing model robustness to rotations.
<b>Sample-wise normalization</b>	A technique to set the mean of each sample to 0 and normalize each sample to have a standard deviation of 1.
<b>Semantic segmentation</b>	A deep learning task that involves classifying each pixel in an image into a predefined class.
<b>Shear range</b>	A data augmentation parameter that applies a shear transformation to an image, slanting it along one axis to simulate different perspectives.
<b>Stride</b>	A parameter in convolution that determines the step size of the kernel when moving across the input data.
<b>TensorFlow</b>	An open-source machine learning library used for various tasks, including deep learning and image processing.

Term	Definition
<b>Transfer learning</b>	A method where a pre-trained model is adapted to a new, related task by adjusting its weights, allowing it to perform well even with limited data for the new task.
<b>Transpose convolution</b>	An operation that reverses the effects of convolution, often used for up-sampling in image processing.
<b>VGG16</b>	A convolutional neural network model pre-trained on the ImageNet data set, commonly used in transfer learning for tasks involving image classification.
<b>Width shift range</b>	A data augmentation parameter that randomly shifts an image horizontally, altering its position to improve model robustness to horizontal translations.
<b>Zoom range</b>	A data augmentation parameter that randomly zooms in or out on an image, altering its scale during training.